

This probably didn't include all the required profiling/index results, but the other profiling/index results can be found in the previous attempt

Assignment 7

Task 1: Normal forms

1. Is this schema in 1NF?
 - a. A schema is 1NF if and only if all underlying domains have atomic values.
All data is atomic if we assume "timestamp" is atomic.
2. Is it in 2NF?
 - a. A schema is 2NF if and only if it is 1NF and all non-key attributes are associated with a candidate key or another non-key attribute. It can be argued that "Passing" does not follow this rule, as "TollStation" does not depend on either "regno" or "timestamp".
3. Is it in 3NF?
 - a. If you assume "station" is linked to "timestamp" and "regno", it is.
Otherwise, it is not since it does not fill 2NF.

This probably didn't include all the required profiling/index results, but the other profiling/index results can be found in the previous attempt

Task 2: Create a physical schema and a test environment

Installed mysqltuner

Sudo dnf install mysqltuner -y

Signed into mariadb

Mariadb -u odnerindheim -p

Created database and use it.

CREATE DATABASE car_db;

USE car_db

Creating aux table for data

CREATE TABLE `temp` (

```
`passing_datetime` datetime DEFAULT NULL,  
`regno` char(7) DEFAULT NULL,  
`tollstation_id` smallint(5) unsigned DEFAULT NULL,  
`tollstation_name` varchar(85) DEFAULT NULL,  
`owner_name` varchar(85) DEFAULT NULL,  
`taxclass_id` smallint(5) unsigned DEFAULT NULL,  
`taxclass_description` text DEFAULT NULL,  
`has_subscription` enum('yes','no') DEFAULT NULL,  
`fee` decimal(5,2) DEFAULT NULL
```

) ENGINE=InnoDB DEFAULT CHARSET=utf8 COLLATE=utf8_general_ci

Downloading the data file:

Wget <https://eple.hvl.no/fag/dat151/v2024/carpassingdb.txt>

Load the data file:

Used DBEAVER to load all the data.

Run some queries:

SELECT COUNT(*) FROM car_passing aux;

```
ERROR 1146 (42S02): Table 'car_db.car_passing'  
[MariaDB [car_db]]> SELECT COUNT(*) FROM temp;  
+-----+  
| COUNT(*) |  
+-----+  
| 4154196 |  
+-----+  
1 row in set (0.991 sec)
```

SELECT * FROM car_passing_aux ORDER BY passing_timestamp ASC LIMIT 20;

Odne Rindheim

Redone Task 2 and Task 3

This probably didn't include all the required profiling/index results, but the other profiling/index results can be found in the previous attempt

passing_datetime	regno	tollstation_id	tollstation_name	owner_name	taxclass_id	taxclass_description	has_subscription	fee
2022-08-16 03:48:04	SR51455	19	Strandkaien	Magnar Kristiansen	2	Ladbar hybrid, taxgroup 1	yes	20.00
2022-08-19 20:35:04	SR51455	21	Sotraveien	Magnar Kristiansen	2	Ladbar hybrid, taxgroup 1	yes	20.00
2022-08-23 23:07:04	SR51455	7	Straume bro	Magnar Kristiansen	2	Ladbar hybrid, taxgroup 1	yes	20.00
2022-08-29 01:49:04	SR51455	23	Straumeveien	Magnar Kristiansen	2	Ladbar hybrid, taxgroup 1	yes	20.00
2022-08-31 16:02:59	SR34676	24	Hardangervegen	Agathe Sæbø	3	Ladbar hybrid, taxgroup 2	no	37.00
2022-09-04 00:22:04	SR51455	13	Michael Krohnsgate	Magnar Kristiansen	2	Ladbar hybrid, taxgroup 1	yes	20.00
2022-09-04 19:17:59	SR34676	23	Straumeveien	Agathe Sæbø	3	Ladbar hybrid, taxgroup 2	no	37.00
2022-09-08 11:01:04	SR51455	28	Flyplassvegen	Magnar Kristiansen	2	Ladbar hybrid, taxgroup 1	yes	20.00
2022-09-09 05:23:59	SR34676	17	Åsamyrene	Agathe Sæbø	3	Ladbar hybrid, taxgroup 2	no	37.00
2022-09-09 21:43:46	EV65954	17	Åsamyrene	Ekaterina Ismail	5	Elbil, taxgroup 1	yes	8.00
2022-09-10 04:40:46	EV65954	10	Florida	Ekaterina Ismail	5	Elbil, taxgroup 1	yes	8.00
2022-09-12 09:45:04	SR51455	15	Tellevikvegen	Magnar Kristiansen	2	Ladbar hybrid, taxgroup 1	yes	20.00
2022-09-13 14:48:04	SR51455	11	Gyldepris	Magnar Kristiansen	2	Ladbar hybrid, taxgroup 1	yes	20.00
2022-09-14 00:38:17	EB36167	23	Straumeveien	Yumiko Hvidsten	5	Elbil, taxgroup 1	no	10.00
2022-09-14 05:26:22	SP61094	23	Straumeveien	Sandra Ødegård	2	Ladbar hybrid, taxgroup 1	no	25.00
2022-09-14 07:58:17	EB36167	25	Fritz C. Riebers vei	Yumiko Hvidsten	5	Elbil, taxgroup 1	no	10.00
2022-09-14 08:17:59	SR34676	10	Florida	Agathe Sæbø	3	Ladbar hybrid, taxgroup 2	no	37.00
2022-09-14 12:34:59	SR34676	12	Kalfaret	Agathe Sæbø	3	Ladbar hybrid, taxgroup 2	no	37.00
2022-09-14 12:46:22	SP61094	25	Fritz C. Riebers vei	Sandra Ødegård	2	Ladbar hybrid, taxgroup 1	no	25.00
2022-09-14 13:42:46	EV65954	29	Skagevegen	Ekaterina Ismail	5	Elbil, taxgroup 1	yes	8.00

SELECT * FROM car_passing_aux ORDER BY passing_timestamp DESC LIMIT 20;

passing_datetime	regno	tollstation_id	tollstation_name	owner_name	taxclass_id	taxclass_description	has_subscription	fee
2024-02-08 20:09:47	EV45094	19	Strandkaien	Stephanie Ravndal	5	Elbil, taxgroup 1	yes	8.00
2024-02-08 20:05:39	SR58844	15	Tellevikvegen	Xhevrije Hoem	4	Diesel, taxgroup 1	yes	24.00
2024-02-08 20:00:19	SR24250	13	Michael Krohnsgate	Lillian Holen	4	Diesel, taxgroup 1	no	30.00
2024-02-08 19:56:12	EV19636	12	Kalfaret	Zhian Osmundsen	5	Elbil, taxgroup 1	yes	8.00
2024-02-08 19:56:08	SP65600	28	Flyplassvegen	Juni Gode	1	Bensin, taxgroup 1	no	25.00
2024-02-08 19:54:23	EB27313	6	Gravdal	Truls Fjellheim	5	Elbil, taxgroup 1	yes	8.00
2024-02-08 19:53:04	ED57003	24	Hardangervegen	Ted Moen	5	Elbil, taxgroup 1	yes	8.00
2024-02-08 19:52:59	SP58784	14	Damsgårdsvingen	Q1 Fremstad	2	Ladbar hybrid, taxgroup 1	no	25.00
2024-02-08 19:49:11	EF50375	18	Åsaneveien	Marlene Sævik	5	Elbil, taxgroup 1	no	10.00
2024-02-08 19:46:45	EB93110	16	Arnavegen	Joar Øie	6	Elbil, taxgroup 2	yes	0.00
2024-02-08 19:46:21	EV87230	20	Småstrandgaten	Wojciech Henningsen	5	Elbil, taxgroup 1	no	10.00
2024-02-08 19:45:28	EV36800	11	Gyldepris	Yenny Kallevik	5	Elbil, taxgroup 1	yes	8.00
2024-02-08 19:45:03	ECB7982	1	Fjøsangerveien	Queen Steinsvik	5	Elbil, taxgroup 1	no	10.00
2024-02-08 19:42:20	SX54186	4	Dolviken	Zanib Grimstad	2	Ladbar hybrid, taxgroup 1	no	25.00
2024-02-08 19:41:40	EB68077	29	Skagevegen	Eilev Jonassen	5	Elbil, taxgroup 1	yes	8.00
2024-02-08 19:38:58	EBG2673	19	Strandkaien	Ulrica Voll	5	Elbil, taxgroup 1	no	10.00
2024-02-08 19:38:31	SP58647	10	Florida	Cassandra Førland	2	Ladbar hybrid, taxgroup 1	yes	20.00
2024-02-08 19:37:13	ST42466	7	Straume bro	Valentina Ask	4	Diesel, taxgroup 1	yes	24.00
2024-02-08 19:37:10	SR60544	7	Straume bro	Zora Voll	2	Ladbar hybrid, taxgroup 1	no	25.00
2024-02-08 19:35:42	SP40238	21	Sotraveien	Ouyen Reistad	2	Ladbar hybrid, taxgroup 1	no	25.00

Run the mysql tuner

Mysqltuner

This probably didn't include all the required profiling/index results, but the other profiling/index results can be found in the previous attempt

```
----- InnoDB Metrics -----
[=] InnoDB is enabled.
[=] InnoDB Thread Concurrency: 0
[OK] InnoDB File per table is activated
[!] InnoDB buffer pool / data size: 128.0M/468.0M
[!] Ratio InnoDB log file size / InnoDB Buffer pool size (75 %): 96.0M * 1/128.0M should be equal to 25%
[=] Number of InnoDB Buffer Pool Chunk : 1 for 1 Buffer Pool Instance(s)
[OK] Innodb_buffer_pool_size aligned with Innodb_buffer_pool_chunk_size & Innodb_buffer_pool_instances
[OK] InnoDB Read buffer efficiency: 99.99% (46894525 hits/ 46900422 total)
[OK] InnoDB Write log efficiency: 99.24% (58893 hits/ 59344 total)
[OK] InnoDB log waits: 0.00% (0 waits / 451 writes)

----- Aria Metrics -----
[--] Aria Storage Engine is enabled.
[OK] Aria pagecache size / total Aria indexes: 128.0M/320.0K
[!] Aria pagecache hit rate: 52.0% (25 cached / 12 reads)

----- TokuDB Metrics -----
[--] TokuDB is disabled.

----- XtraDB Metrics -----
[--] XtraDB is disabled.

----- Galera Metrics -----
[--] Galera is disabled.

----- Replication Metrics -----
[--] Galera Synchronous replication: NO
[--] No replication slave(s) for this server.
[--] Binlog format: MIXED
[--] XA support enabled: ON
[--] Semi synchronous replication Master: OFF
[--] Semi synchronous replication Slave: OFF
[--] This is a standalone server

----- Recommendations -----
General recommendations:
  Reduce or eliminate unclosed connections and network issues
  Configure your accounts with ip or subnets only, then update your configuration with skip-name-resolve=1
  Performance schema should be activated for better diagnostics
  Consider installing Sys schema from https://github.com/mysql/mysql-sys for MySQL
  Consider installing Sys schema from https://github.com/FromDual/mariadb-sys for MariaDB
  Before changing innodb_log_file_size and/or innodb_log_files_in_group read this: https://bit.ly/2TcGgtU
Variables to adjust:
  performance_schema = ON enable PFS
  innodb_buffer_pool_size (>= 468.0M) if possible.
  innodb_log_file_size should be (=32M) if possible, so InnoDB total log files size equals to 25% of buffer pool size.
[odnerindheim@localhost ~]$ mysqltuner
>> MySQLTuner 1.8.3 - Major Hayden <major@mhtx.net>
>> Bug reports, feature requests, and downloads at http://mysqltuner.pl/
>> Run with '--help' for additional options and output filtering
```

Increasing the `buffer_pool_size`:

Using `SET GLOBAL` as this doesn't require a restart: [MariaDB Enterprise Server](#)
[Configure the InnoDB Buffer Pool — MariaDB Documentation](#)

```
SET GLOBAL innodb_buffer_pool_size=(2*1024*1024*1024);
```

```
SHOW GLOBAL STATUS
```

```
    LIKE 'Innodb_buffer_pool_resize_status';
```

```
MariaDB [(none)]> SHOW GLOBAL STATUS
  -> LIKE 'Innodb_buffer_pool_resize_status';
+-----+-----+
| Variable_name          | Value           |
+-----+-----+
| Innodb_buffer_pool_resize_status | Completed resizing buffer pool at 240419 16:38:43. |
+-----+-----+
1 row in set (0.000 sec)
```

Create normalized tables:

```
-- car_db.TaxClass definition
```

```
CREATE TABLE `TaxClass` (
  `taxclass` smallint(5) unsigned NOT NULL AUTO_INCREMENT,
  `description` text NOT NULL,
  PRIMARY KEY (`taxclass`)
) ENGINE=InnoDB DEFAULT CHARSET=utf8 COLLATE=utf8_general_ci;
```

This probably didn't include all the required profiling/index results, but the other profiling/index results can be found in the previous attempt

-- car_db.Tollstation definition

```
CREATE TABLE `Tollstation` (
  `tollstation` smallint(5) unsigned NOT NULL AUTO_INCREMENT,
  `name` varchar(85) NOT NULL,
  PRIMARY KEY (`tollstation`)
) ENGINE=InnoDB DEFAULT CHARSET=utf8 COLLATE=utf8_general_ci;
```

-- car_db.Car definition

```
CREATE TABLE `Car` (
  `regno` char(7) NOT NULL,
  `owner` varchar(85) NOT NULL,
  `taxclass` smallint(5) unsigned NOT NULL,
  PRIMARY KEY (`regno`),
  KEY `taxclass` (`taxclass`),
  CONSTRAINT `Car_ibfk_1` FOREIGN KEY (`taxclass`) REFERENCES `TaxClass`(`taxclass`)
) ENGINE=InnoDB DEFAULT CHARSET=utf8 COLLATE=utf8_general_ci;
```

-- car_db.Fee definition

```
CREATE TABLE `Fee` (
  `fee` smallint(5) unsigned NOT NULL AUTO_INCREMENT,
  `taxclass` smallint(5) unsigned NOT NULL,
  `type` enum('regular','withsubscription') DEFAULT NULL,
  `costPerPassing` decimal(5,2) DEFAULT NULL,
  PRIMARY KEY (`fee`),
  UNIQUE KEY `taxclass` (`taxclass`,`type`),
  CONSTRAINT `Fee_ibfk_1` FOREIGN KEY (`taxclass`) REFERENCES `TaxClass`(`taxclass`)
) ENGINE=InnoDB DEFAULT CHARSET=utf8 COLLATE=utf8_general_ci;
```

-- car_db.Passing definition

```
CREATE TABLE `Passing` (
  `timestamp` datetime NOT NULL,
  `regno` char(7) NOT NULL,
  `tollstation` smallint(5) unsigned NOT NULL,
  PRIMARY KEY (`timestamp`, `regno`),
  KEY `tollstation` (`tollstation`),
  KEY `regno` (`regno`),
```

This probably didn't include all the required profiling/index results, but the other profiling/index results can be found in the previous attempt

```
CONSTRAINT `Passing_ibfk_1` FOREIGN KEY (`tollstation`) REFERENCES
`Tollstation`(`tollstation`),
CONSTRAINT `Passing_ibfk_2` FOREIGN KEY (`regno`) REFERENCES `Car`
(`regno`)
) ENGINE=InnoDB DEFAULT CHARSET=utf8 COLLATE=utf8_general_ci;
```

-- car_db.Subscription definition

```
CREATE TABLE `Subscription` (
`regno` char(7) NOT NULL,
PRIMARY KEY (`regno`),
CONSTRAINT `Subscription_ibfk_1` FOREIGN KEY (`regno`) REFERENCES `Car`
(`regno`)
) ENGINE=InnoDB DEFAULT CHARSET=utf8 COLLATE=utf8_general_ci;
```

Populate tables:

Insert into Tollstation (tollstation, name) select distinct tollstation_id, tollstation_name from temp;

Insert into TaxClass (taxclass, desciption) select distinct taxclass_id, taxclass_description from temp;

Insert into Car (regno, owner, taxclass) select distinct regno, owner_name, taxclass_id from temp;

Insert into Subscription select distinct regno from temp where has_subscription='yes';

Insert into passing (timestamp,regno,tollstation) select distinct passing_datetime, regno, tollstation_id from temp;

Insert into fee(taxclass,type,costperpassing) select distinct taxclass_id,'regular',fee from temp where has_subscription='no';

Insert into fee(taxclass,type,costperpassing) select distinct taxclass_id,'withsubscription',fee from temp where has_subscription='yes';

This probably didn't include all the required profiling/index results, but the other profiling/index results can be found in the previous attempt

Verifying data integrity:

```
SELECT * FROM Tollstation LIMIT 10;  
SELECT * FROM TaxClass LIMIT 10;  
SELECT * FROM Car LIMIT 10;  
SELECT * FROM Subscription LIMIT 10;  
SELECT * FROM Passing LIMIT 10;  
SELECT * FROM Fee LIMIT 10;
```

All data looks in order!

Dropping temp

```
DROP TABLE temp;
```

This probably didn't include all the required profiling/index results, but the other profiling/index results can be found in the previous attempt

Task 3: Performance of the physical model

Cache was cleared and mariadb was restarted before each query.

Query 1:

A)

Finds the owner name of passings of a certain time and date.

```
SELECT C.owner, P.timestamp
FROM Car C JOIN Passing P USING(regno)
WHERE YEAR(P.timestamp)=2023 AND MONTH(P.timestamp)=2
AND DAYOFMONTH(P.timestamp)=1 AND HOUR(P.timestamp)=23;
```

B)

SET profiling = 1;

Running the query 3 times.

```
MariaDB [car_db]> SHOW PROFILES;
```

Query_ID	Duration	Query
1 1.02631599		SELECT C.owner, P.timestamp FROM Car C JOIN Passing P USING (regno) WHERE YEAR(P.timestamp)=2023 AND MONTH(P.timestamp)=2 AND DAYOFMONTH(P.timestamp)=1 AND HOUR(P.timestamp)=23
2 1.01596361		SELECT C.owner, P.timestamp FROM Car C JOIN Passing P USING (regno) WHERE YEAR(P.timestamp)=2023 AND MONTH(P.timestamp)=2 AND DAYOFMONTH(P.timestamp)=1 AND HOUR(P.timestamp)=23
3 1.02022030		SELECT C.owner, P.timestamp FROM Car C JOIN Passing P USING (regno) WHERE YEAR(P.timestamp)=2023 AND MONTH(P.timestamp)=2 AND DAYOFMONTH(P.timestamp)=1 AND HOUR(P.timestamp)=23

```
MariaDB [car_db]> EXPLAIN SELECT C.owner, P.timestamp FROM Car C JOIN Passing P USING (regno) WHERE
YEAR(P.timestamp)=2023 AND MONTH(P.timestamp)=2 AND DAYOFMONTH(P.timestamp)=1 AND HOUR(P.timestamp)=23;
+-----+-----+-----+-----+-----+-----+-----+
| id | select_type | table | type | possible_keys | key | key_len | ref      | rows | Extra          |
+-----+-----+-----+-----+-----+-----+-----+
| 1 | SIMPLE    | C    | PRIMARY | NULL    | NULL    | NULL    | NULL     | 62896 |                |
| 1 | SIMPLE    | P    | ref     | regno   | 17     | car_db.C.regno | Using where; Using index |
+-----+-----+-----+-----+-----+-----+-----+
```

2 rows in set (0.000 sec)

C)

Denormalising Passing:

I made an attempt to denormalize passing table, to avoid using function calls like YEAR(), MONTH(), DAYOFMONTH(), HOUR(), which can prevent the use of index:

CREATE TABLE Passing_denorm AS

SELECT P*,

```
YEAR(P.timestamp) AS year,
MONTH(P.timestamp) AS month,
DAYOFMONTH(P.timestamp) AS day,
HOUR(P.timestamp) AS hour
```

This probably didn't include all the required profiling/index results, but the other profiling/index results can be found in the previous attempt

FROM Passing P;

new query:

```
SELECT C.owner, P.timestamp
FROM Car C JOIN Passing_denorm P USING(regno)
WHERE P.year = 2023 AND P.month = 2
AND P.day = 1 AND P.hour = 23;
```

Profiling results (mariadb was restarted and cache was cleared):

```
MariaDB [car_db]> SHOW PROFILES;
+-----+-----+
--+
| Query_ID | Duration | Query
+-----+-----+
--+
| 1 | 1.35718505 | SELECT C.owner, P.timestamp FROM Car C JOIN Passing_denorm P USING(regno) WHERE P.year = 2023 AND P.month = 2 AND P.day = 1 AND P.hour = 23 |
| 2 | 1.36653832 | SELECT C.owner, P.timestamp FROM Car C JOIN Passing_denorm P USING(regno) WHERE P.year = 2023 AND P.month = 2 AND P.day = 1 AND P.hour = 23 |
| 3 | 1.36108667 | SELECT C.owner, P.timestamp FROM Car C JOIN Passing_denorm P USING(regno) WHERE P.year = 2023 AND P.month = 2 AND P.day = 1 AND P.hour = 23 |
+-----+-----+
--+
3 rows in set (0.000 sec)
```

Not exactly better results, trying explain:

```
MariaDB [car_db]> EXPLAIN SELECT C.owner, P.timestamp FROM Car C JOIN Passing_denorm P USING(regno) WHERE P.year = 2023 AND P.month = 2 AND P.day = 1 AND P.hour = 23;
+-----+-----+-----+-----+-----+-----+-----+-----+
| id | select_type | table | type | possible_keys | key | key_len | ref | rows | Extra |
+-----+-----+-----+-----+-----+-----+-----+-----+
| 1 | SIMPLE | P | ALL | NULL | NULL | NULL | 4138561 | Using where |
| 1 | SIMPLE | C | eq_ref | PRIMARY | PRIMARY | 17 | car_db.P.regno | 1 | |
+-----+-----+-----+-----+-----+-----+-----+-----+
2 rows in set (0.000 sec)
```

This denormalization removed the use of index.

This denormalization violates the 3NF as it introduces transitive dependencies – each one of these date parts is non-prime and is determined by ‘timestamp’, which is not part of the primary key.

I can try denormalizing by adding a owner column to passing:

This denormalization also violates normalization rule of storing non-key attributes (attributes that do not contribute to a composite primary key), in this case ‘owner’, which is related to car and not directly to the passing event.

```
CREATE TABLE Passing_with_owner AS
SELECT P.timestamp, P.regno, P.tollstation, C.owner
FROM Passing P
JOIN Car C ON P.regno = C.regno;
```

Optimized query:

```
SELECT owner, timestamp
FROM Passing_with_owner
WHERE YEAR(timestamp) = 2023 AND MONTH(timestamp) = 2
```

Odne Rindheim

Redone Task 2 and Task 3

This probably didn't include all the required profiling/index results, but the other profiling/index results can be found in the previous attempt

AND DAYOFMONTH(timestamp) = 1 AND HOUR(timestamp) = 23;

Profile:

MariaDB [car_db]> SHOW PROFILES;

```
+-----+-----+
|-----+
| Query_ID | Duration | Query
|-----+-----+
|-----+
|   1 | 1.33344789 | SELECT owner,timestamp FROM Passing_with_owner WHERE YEAR(timestamp)=2023 AND MONTH(timestamp)=2 AND DAYOFMONTH(timestamp)=1 AND HOUR(timestamp)=23 |
|   2 | 1.32670961 | SELECT owner,timestamp FROM Passing_with_owner WHERE YEAR(timestamp)=2023 AND MONTH(timestamp)=2 AND DAYOFMONTH(timestamp)=1 AND HOUR(timestamp)=23 |
|   3 | 1.32485803 | SELECT owner,timestamp FROM Passing_with_owner WHERE YEAR(timestamp)=2023 AND MONTH(timestamp)=2 AND DAYOFMONTH(timestamp)=1 AND HOUR(timestamp)=23 |
+-----+-----+
|-----+
3 rows in set (0.000 sec)
```

```
MariaDB [car_db]> EXPLAIN SELECT owner,timestamp FROM Passing_with_owner WHERE YEAR(timestamp)=2023 AND MONTH(timestamp)=2 AND DAYOFMONTH(timestamp)=1 AND HOUR(timestamp)=23;
```

D)

Could I make the dbms use index in this query? To make the dbms more likely to use an index, I could change the query to use a range-based function instead of functions, heres how I could create an index, and use it:

```
CREATE INDEX idx_passing_owner_timestamp ON Passing WITH OWNER(timestamp);
```

```
SELECT owner, timestamp  
FROM Passing_with_owner  
WHERE timestamp >= '202
```

```
MariaDB [car_db]> EXPLAIN SELECT owner, timestamp FROM Passing_with_owner WHERE timestamp >= '2023-02-01 23:00:00'
```

```
;;
+-----+-----+-----+-----+-----+-----+-----+-----+
| id | select_type | table      | type   | possible_keys    | key           | key_len | ref | rows | Extra       |
+-----+-----+-----+-----+-----+-----+-----+-----+
| 1 | SIMPLE    | Passing_with_owner | range | idx_passing_owner_timestamp | idx_passing_owner_timestamp | 5 | NULL | 151 | Using index condition |
+-----+-----+-----+-----+-----+-----+-----+-----+
1 row in set (0.001 sec)
```

Creating an index made the Select function very fast:

```
CREATE INDEX idx_passing_with_owner ON Passing_with_owner(timestamp, owner);
```

```
MariaDB [car_db]> EXPLAIN SELECT owner,timestamp FROM Passing_with_owner WHERE timestamp >= '2023-02-01 23:00:00' AND timestamp <'2023-02-02 00:00:00';
```

```
+-----+-----+-----+-----+-----+-----+-----+-----+
| id | select_type | table      | type | possible_keys          | key      | key_len | ref | rows | Extra      |
+-----+-----+-----+-----+-----+-----+-----+-----+
```

Odne Rindheim

Redone Task 2 and Task 3

This probably didn't include all the required profiling/index results, but the other profiling/index results can be found in the previous attempt

```
+-----+-----+-----+-----+-----+-----+
-----+
| 1 | SIMPLE  | Passing_with_owner | range | idx_passing_owner_timestamp, idx_passing_with_owner | idx_passing_with_owner | 5
| NULL | 151 | Using where; Using index |
+-----+-----+-----+-----+-----+-----+
-----+
1 row in set (0.000 sec)
MariaDB [car_db]> SHOW PROFILES;
+-----+-----+-----+
| Query_ID | Duration | Query
+-----+-----+-----+
| 1 | 0.00244040 | SELECT owner,timestamp FROM Passing_with_owner WHERE timestamp >= '2023-02-01 23:00:00' AND
timestamp <'2023-02-02 00:00:00' |
| 2 | 0.00038712 | SELECT owner,timestamp FROM Passing_with_owner WHERE timestamp >= '2023-02-01 23:00:00' AND
timestamp <'2023-02-02 00:00:00' |
| 3 | 0.00038903 | SELECT owner,timestamp FROM Passing_with_owner WHERE timestamp >= '2023-02-01 23:00:00' AND
timestamp <'2023-02-02 00:00:00' |
+-----+-----+-----+
3 rows in set (0.000 sec)
```

This probably didn't include all the required profiling/index results, but the other profiling/index results can be found in the previous attempt

Query 2:

- A) Retrieves the registration numbers of cars that have a subscription and calculates the total fee for each car, where the total fee is more than 5000.
- B) Profiling results:

```
MariaDB [car_db]> show profiles;
+-----+-----+
| Query_ID | Duration   | Query
+-----+-----+
|      18 | 6.41802076 | SELECT C.regno AS regno, Sum(F.costPerPassing) AS totalfee FROM Car C JOIN Passing P USING(regno) JOIN TaxClass T USING(taxclass)
N Fee F USING(taxclass) WHERE F.type='withsubscription' AND C.regno IN (SELECT regno FROM Subscription) GROUP BY C.regno HAVING totalfee > 5000 |
|      19 | 6.42284271 | SELECT C.regno AS regno, Sum(F.costPerPassing) AS totalfee FROM Car C JOIN Passing P USING(regno) JOIN TaxClass T USING(taxclass)
N Fee F USING(taxclass) WHERE F.type='withsubscription' AND C.regno IN (SELECT regno FROM Subscription) GROUP BY C.regno HAVING totalfee > 5000 |
|      20 | 6.32820439 | SELECT C.regno AS regno, Sum(F.costPerPassing) AS totalfee FROM Car C JOIN Passing P USING(regno) JOIN TaxClass T USING(taxclass)
N Fee F USING(taxclass) WHERE F.type='withsubscription' AND C.regno IN (SELECT regno FROM Subscription) GROUP BY C.regno HAVING totalfee > 5000 |
+-----+-----+
3 rows in set (0.000 sec)
```

- C) Can create a table that stores the pre-calculated sums of fees for cars with subscriptions. This way, we avoid the need to calculate the sum dynamically each time the query is executed. Since both attributes are atomic values, and there are no partial dependencies it is 1NF and 2NF, but since totalfee attribute is transitively dependent on the car tables attributes, it violates 3NF

Create Denormalized Table:

```
CREATE TABLE CarTotalFees (
regno CHAR(7) PRIMARY KEY,
totalFee DECIMAL(10,2)
);
```

```
INSERT INTO CarTotalFees (regno, totalFee)
SELECT C.regno, SUM(F.costPerPassing) AS totalfee
FROM Car C
JOIN Passing P USING (regno)
JOIN Fee F USING (taxclass)
JOIN Subscription S USING (regno)
WHERE F.type = 'withsubscription'
GROUP BY C.regno;
```

```
SELECT *
FROM CarTotalFees
WHERE totalFee > 5000;
```

Profiling results:

```
[MariaDB [car_db]> show profiles;
+-----+-----+
| Query_ID | Duration   | Query
+-----+-----+
|      23 | 0.00665651 | SELECT * FROM CarTotalFees WHERE totalFee > 5000 |
|      24 | 0.00669559 | SELECT * FROM CarTotalFees WHERE totalFee > 5000 |
|      25 | 0.00668964 | SELECT * FROM CarTotalFees WHERE totalFee > 5000 |
+-----+-----+
3 rows in set (0.000 sec)
```

This probably didn't include all the required profiling/index results, but the other profiling/index results can be found in the previous attempt

D) Explain the original query:

```
MariaDB [car_db]> EXPLAIN SELECT C.regno AS regno, Sum(F.costPerPassing) AS totalfee FROM Car C JOIN Passing P USING(regno) JOIN TaxClass T USING(taxclass) JOIN Fee F USING(taxclass)
+-----+-----+-----+-----+-----+-----+-----+-----+
| id | select_type | table | type | possible_keys | key | key_len | ref | rows | Extra
+-----+-----+-----+-----+-----+-----+-----+-----+
| 1 | PRIMARY | T | ALL | PRIMARY | NULL | NULL | NULL | 10 | Using temporary; Using filesort
| 1 | PRIMARY | F | eq_ref | taxclass | taxclass | 4 | car_db.T.taxclass,const | 1 | Using index condition
| 1 | PRIMARY | C | ref | PRIMARY,taxclass | taxclass | 2 | car_db.T.taxclass | 1 | Using index
| 1 | PRIMARY | Subscription | eq_ref | PRIMARY | PRIMARY | 21 | car_db.C.regno | 1 | Using index
| 1 | PRIMARY | P | ref | regno | regno | 21 | car_db.C.regno | 1 | Using index
+-----+-----+-----+-----+-----+-----+-----+-----+
```

E) Created some indexes, to see if it could speed up original quieuyes:

CREATE INDEX idx_regno ON Car (regno);

CREATE INDEX idx_type ON Fee (type, taxclass, costPerPassing);

CREATE INDEX idx_regno_subscription ON Subscription (regno);

```
MariaDB [car_db]> SHOW PROFILES;
+-----+-----+-----+
| Query_ID | Duration | Query
+-----+-----+-----+
| 28 | 6.31994821 | SELECT C.regno AS regno, Sum(F.costPerPassing) AS totalfee FROM Car C JOIN Passing P USING(regno) JOIN TaxClass T USING(taxclass) JOIN Fee F USING(taxclass) WHERE F.type='withsubscription' AND C.regno IN (SELECT regno FROM Subscription) GROUP BY C.regno HAVING totalfee > 5000 |
| 29 | 6.31393725 | SELECT C.regno AS regno, Sum(F.costPerPassing) AS totalfee FROM Car C JOIN Passing P USING(regno) JOIN TaxClass T USING(taxclass) JOIN Fee F USING(taxclass) WHERE F.type='withsubscription' AND C.regno IN (SELECT regno FROM Subscription) GROUP BY C.regno HAVING totalfee > 5000 |
| 30 | 6.32113335 | SELECT C.regno AS regno, Sum(F.costPerPassing) AS totalfee FROM Car C JOIN Passing P USING(regno) JOIN TaxClass T USING(taxclass) JOIN Fee F USING(taxclass) WHERE F.type='withsubscription' AND C.regno IN (SELECT regno FROM Subscription) GROUP BY C.regno HAVING totalfee > 5000 |
+-----+-----+-----+
MariaDB [car_db]> EXPLAIN SELECT C.regno AS regno, Sum(F.costPerPassing) AS totalfee FROM Car C JOIN Passing P USING(regno) JOIN TaxClass T USING(taxclass) JOIN Fee F USING(taxclass) WHERE F.type='withsubscription' AND C.regno IN (SELECT regno FROM Subscription) GROUP BY C.regno HAVING totalfee > 5000;
+-----+-----+-----+-----+-----+-----+-----+-----+
| id | select_type | table | type | possible_keys | key | key_len | ref | rows | Extra
+-----+-----+-----+-----+-----+-----+-----+-----+
| 1 | PRIMARY | F | ref | taxclass,IDX_TYPE | IDX_TYPE | 2 | const | 10 | Using where; Using index; Using temporary; Using filesort
| 1 | PRIMARY | T | eq_ref | PRIMARY | PRIMARY | 2 | car_db.F.taxclass | 1 |
| 1 | PRIMARY | C | ref | PRIMARY,taxclass,IDX_REGNO | taxclass | 2 | car_db.F.taxclass | 3593 | Using index
| 1 | PRIMARY | Subscription | eq_ref | PRIMARY,IDX_REGNO_SUBSCRIPTION | PRIMARY | 21 | car_db.C.regno | 1 | Using index
| 1 | PRIMARY | P | ref | regno | regno | 21 | car_db.C.regno | 1 |
+-----+-----+-----+-----+-----+-----+-----+-----+
5 rows in set (0.001 sec)
```

This probably didn't include all the required profiling/index results, but the other profiling/index results can be found in the previous attempt

Query 3

- A) Like query two, but instead of directly joining with the subscription table, it uses a subquery to find cars with subscriptions.
 - B) Profiling results:

```
MariaDB [car_db]> show profiles;
+-----+-----+-----+
| Query_ID | Duration | Query
+-----+-----+-----+
| 1 | 35 | 6.29570014 | SELECT C.regno AS regno, Sum(F.costPerPassing) AS totalfee FROM Car C JOIN Passing P USING(regno) JOIN TaxClass T USING(taxclass) JOIN Fee F USING(taxclass) WHERE F.type='withsubscription' AND C.regno IN (SELECT regno FROM Subscription) GROUP BY C.regno HAVING totalfee > 5000 |
| 36 | 6.29970013 | SELECT C.regno AS regno, Sum(F.costPerPassing) AS totalfee FROM Car C JOIN Passing P USING(regno) JOIN TaxClass T USING(taxclass) JOIN Fee F USING(taxclass) WHERE F.type='withsubscription' AND C.regno IN (SELECT regno FROM Subscription) GROUP BY C.regno HAVING totalfee > 5000 |
| 37 | 6.29809210 | SELECT C.regno AS regno, Sum(F.costPerPassing) AS totalfee FROM Car C JOIN Passing P USING(regno) JOIN TaxClass T USING(taxclass) JOIN Fee F USING(taxclass) WHERE F.type='withsubscription' AND C.regno IN (SELECT regno FROM Subscription) GROUP BY C.regno HAVING totalfee > 5000 |
+-----+-----+-----+
3 rows in set (0.000 sec)
```

- C) Possible denormalizing the query, very similar to the denormalizing of query 2, this also breaks the 3NF by introducing redundancy.

CREATE TABLE query 3 (

regno CHAR(7) PRIMARY KEY,
totalFee DECIMAL (10,2)

1

Inserting data into new table:

INSERT INTO query3 (regno, totalFee)

```
-> SELECT C.regno AS regno, SUM(F.costPerPassing) AS totalfee  
-> FROM Car C JOIN Passing P USING (regno)  
-> JOIN TaxClass T USING(taxclass)  
-> JOIN Fee F USING(taxclass)  
-> WHERE F.type='withsubscription'  
-> AND C.regno IN (SELECT regno FROM Subscription)  
-> GROUP BY C.regno;
```

New query:

SELECT *

FROM query3

```
WHERE totalFee > 5000;
```

Profiling results:

```
[MariaDB [car_db]]> show profiles;
+-----+-----+-----+
| Query_ID | Duration | Query
+-----+-----+-----+
| 46 | 0.00679459 | SELECT
| 47 | 0.00669449 | SELECT
| 48 | 0.00678845 | SELECT
+-----+-----+-----+
3 rows in set (0.000 sec)
```

- D) Adding an index to optimize the join and where condition: CREATE INDEX idx_car_with_subscription ON Car_with_subscription(regno, taxclass, subscription);
MariaDB [car_db]> SHOW PROFILES; +-----+-----+

Odne Rindheim

Redone Task 2 and Task 3

This probably didn't include all the required profiling/index results, but the other profiling/index results can be found in the previous attempt

Query_ID | Duration | Query || 1 | 0.74527054 | SELECT CWS.regno,
SUM(F.costPerPassing) AS totalfee FROM Car_with_subscription CWS JOIN Passing P
ON CWS.regno = P.regno JOIN TaxClass T on CWS.taxclass = T.taxclass JOIN Fee F ON
T.taxclass = F.taxclass WHERE F.type = 'withsubscription' AND CWS.subscription = 'yes'
GROUP BY CWS.regno HAVING totalfe || 2 | 0.72084794 | SELECT CWS.regno,
SUM(F.costPerPassing) AS totalfee FROM Car_with_subscription CWS JOIN Passing P
ON CWS.regno = P.regno JOIN TaxClass T on CWS.taxclass = T.taxclass JOIN Fee F ON
T.taxclass = F.taxclass WHERE F.type = 'withsubscription' AND CWS.subscription = 'yes'
GROUP BY CWS.regno HAVING totalfe || 3 | 0.71515412 | SELECT CWS.regno,
SUM(F.costPerPassing) AS totalfee FROM Car_with_subscription CWS JOIN Passing P
ON CWS.regno = P.regno JOIN TaxClass T on CWS.taxclass = T.taxclass JOIN Fee F ON
T.taxclass = F.taxclass WHERE F.type = 'withsubscription' AND CWS.subscription = 'yes'
GROUP BY CWS.regno HAVING totalfe | +-----+-----+

E)

Replacing the ‘IN’ subquery with an ‘INNER JOIN’, which is often more efficient as it can utilize indexes better than a subquery that might result in a full scan. Odne Rindheim This query will allow the DBMS to use indexes efficiently during the join operation instead of possibly running a separate subquery for each row processed in the main query:

```
SELECT C.regno AS regno, SUM(F.costPerPassing) AS totalfee FROM Car C JOIN Passing P USING(regno) JOIN TaxClass T USING(taxclass) JOIN Fee F USING(taxclass) JOIN Subscription S ON C.regno = S.regno WHERE F.type = 'withsubscription' GROUP BY C.regno HAVING totalfee > 5000;
```

Odne Rindheim

Redone Task 2 and Task 3

This probably didn't include all the required profiling/index results, but the other profiling/index results can be found in the previous attempt

This probably didn't include all the required profiling/index results, but the other profiling/index results can be found in the previous attempt

Query 4:

- A) Similar to query two, but it asks for cars without a subscription.

Query:

```
SELECT C.regno AS regno,
       SUM(F.costPerPassing) AS totalfee
  FROM Car C
 JOIN Passing P USING(regno)
 JOIN TaxClass T USING(taxclass)
 JOIN Fee F USING(taxclass)
 LEFT JOIN Subscription S ON C.regno = S.regno
 WHERE F.type = 'withsubscription' AND S.regno IS NULL
 GROUP BY C.regno
 HAVING totalfee > 5000;
```

- B) Profiling results of original query:

```
MariaDB [car_db]> SHOW PROFILES;
+-----+-----+-----+
| Query_ID | Duration | Query
+-----+-----+-----+
|      51 | 4.33079953 | SELECT C.regno AS regno,          SUM(F.costPerPassing) AS totalfee FROM Car C JOIN Passing P USING(regno) JOIN TaxClass T USING(taxclas
s) JOIN Fee F USING(taxclass) LEFT JOIN Subscription S ON C.regno = S.regno WHERE F.type = 'withsubscription' AND S.regno IS NULL GROUP BY C.regno HAVING totalfee
> |
|      52 | 4.34868470 | SELECT C.regno AS regno,          SUM(F.costPerPassing) AS totalfee FROM Car C JOIN Passing P USING(regno) JOIN TaxClass T USING(taxclas
s) JOIN Fee F USING(taxclass) LEFT JOIN Subscription S ON C.regno = S.regno WHERE F.type = 'withsubscription' AND S.regno IS NULL GROUP BY C.regno HAVING totalfee
> |
|      53 | 4.31732821 | SELECT C.regno AS regno,          SUM(F.costPerPassing) AS totalfee FROM Car C JOIN Passing P USING(regno) JOIN TaxClass T USING(taxclas
s) JOIN Fee F USING(taxclass) LEFT JOIN Subscription S ON C.regno = S.regno WHERE F.type = 'withsubscription' AND S.regno IS NULL GROUP BY C.regno HAVING totalfee
> |
+-----+-----+-----+
3 rows in set (0.000 sec)
```

- C) Denormalisation approach would be the same as query 2 and query 3, just with cars without a subscription

Create table query4
 regno CHAR(7) PRIMARY KEY,
 totalFee DECIMAL(10,2)
);

Insert into query4 (regno, totalFee)
 SELECT C.regno, SUM(F.costPerPassing) AS totalfee
 FROM Car C
 JOIN Passing P USING (regno)
 JOIN TaxClass T USING(taxclass)
 JOIN Fee F USING (taxclass)
 LEFT JOIN Subscription S on C.regno = S.regno
 WHERE f.type = 'withsubscription' AND S.regno IS NULL
 GROUP BY C.regno;

SELECT * FRIN query4 WHERE totalFee > 5000;

This probably didn't include all the required profiling/index results, but the other profiling/index results can be found in the previous attempt

Profiling results:

```
MariaDB [car_db]> show profiles;
+-----+-----+
| Query_ID | Duration   | Query
+-----+-----+
|      55 | 0.00460380 | SELECT * FROM CarTotalFeesWithoutSubscription WHERE totalFee > 5000 |
|      56 | 0.00465601 | SELECT * FROM CarTotalFeesWithoutSubscription WHERE totalFee > 5000 |
|      57 | 0.00461748 | SELECT * FROM CarTotalFeesWithoutSubscription WHERE totalFee > 5000 |
+-----+-----+
3 rows in set (0.000 sec)

MariaDB [car_db]> EXPLAIN SELECT C.regno AS regno,          SUM(F.costPerPassing) AS totalfee FROM Car C  JOIN Passing P USING(regno) JOIN TaxClass T USING(taxclas
s) JOIN Fee F USING(taxclass) LEFT JOIN Subscription S ON C.regno = S.regno WHERE F.type = 'withsubscription' AND S.regno IS NULL GROUP BY C.regno HAVING totalfee
> 5000;
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
| id  | select_type | table | type    | possible_keys           | key        | key_len | ref     | rows  | Extra
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
| 1   | SIMPLE      | F    | ref     | taxclass, idx_type     | idx_type   | 2       | const   | 10    | Using where; Using index; Using temporary
| 1   | SIMPLE      | T    | eq_ref   | PRIMARY                | PRIMARY    | 2       | car_db.F.taxclass | 1     |
| 1   | SIMPLE      | C    | ref     | PRIMARY, taxclass, idx_regno | PRIMARY   | 2       | car_db.F.taxclass | 3593 | Using index
| 1   | SIMPLE      | S    | eq_ref   | PRIMARY, idx_regno_subscription | PRIMARY | 21      | car_db.C.regno  | 1     | Using where; Using index; Not exists
| 1   | SIMPLE      | P    | ref     | regno                  | regno     | 21      | car_db.C.regno  | 1     |
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
```

This probably didn't include all the required profiling/index results, but the other profiling/index results can be found in the previous attempt

Query 5:

- a) Similar to query three, but it asks for cars without subscription.
- b) Profiling results of query and query:

```
SELECT C.regno AS regno,      SUM(F.costPerPassing) AS totalfee FROM Car C JOIN
Passing P USING(regno) JOIN TaxClass T USING(taxclass) JOIN
Fee F USING(taxclass) LEFT JOIN Subscription S ON C.regno = S.regno WHERE
F.type = 'withsubscription' AND S.regno IS NULL GROUP BY C.regno HAVING
totalfee > 5000;
```

```
MariaDB [car_db]> show profiles;
+-----+-----+
| Query_ID | Duration |
|-----+-----+
| 59 | 4.23746201 | SELECT C.regno AS regno,      SUM(F.costPerPassing) AS totalfee FROM Car C JOIN Passing P USING(regno) JOIN TaxClass T USING(taxclass) JOIN Fee F USING(taxclass) LEFT JOIN Subscription S ON C.regno = S.regno WHERE F.type = 'withsubscription' AND S.regno IS NULL GROUP BY C.regno HAVING totalfee
> |
| 60 | 4.23484727 | SELECT C.regno AS regno,      SUM(F.costPerPassing) AS totalfee FROM Car C JOIN Passing P USING(regno) JOIN TaxClass T USING(taxclass) JOIN Fee F USING(taxclass) LEFT JOIN Subscription S ON C.regno = S.regno WHERE F.type = 'withsubscription' AND S.regno IS NULL GROUP BY C.regno HAVING totalfee
> |
| 61 | 4.23186782 | SELECT C.regno AS regno,      SUM(F.costPerPassing) AS totalfee FROM Car C JOIN Passing P USING(regno) JOIN TaxClass T USING(taxclass) JOIN Fee F USING(taxclass) LEFT JOIN Subscription S ON C.regno = S.regno WHERE F.type = 'withsubscription' AND S.regno IS NULL GROUP BY C.regno HAVING totalfee
> |
+-----+-----+
3 rows in set (0.000 sec)
```

- c) Denormalizing would be about the same as all the other queries. Creating a table which pre calculates the sum of each of the cars, makes querying the denormalized table simple. But it does compromise on the normalizing of the table and breaks 3NF.

```
CREATE TABLE query5 (
-> regno CHAR(7) PRIMARY KEY,
-> totalFee DECIMAL (10,2)
-> );
INSERT INTO query5 (regno, totalfee)
-> SELECT C.regno, SUM(F.costPerPassing) as totalfee
-> from Car C
-> join Passing P using(regno)
-> join TaxClass T using(taxclass)
-> join Fee F using(taxclass)
-> left join Subscription S on C.regno = S.regno
-> where F.type = 'withsubscription' AND S.regno is null
-> group by C.regno;
```

```
SELECT * FROM query5 WHERE totalFee > 5000;
```

This probably didn't include all the required profiling/index results, but the other profiling/index results can be found in the previous attempt

Profiling results:

```
| MariaDB [car_db]> show profiles;
+-----+-----+-----+
| Query_ID | Duration | Query
+-----+-----+-----+
|      69 | 0.00467680 | select * from query5 where totalfee > 5000 |
|      70 | 0.00464582 | select * from query5 where totalfee > 5000 |
|      71 | 0.00462014 | select * from query5 where totalfee > 5000 |
+-----+-----+-----+
3 rows in set (0.000 sec)

MariaDB [car_db]> EXPLAIN SELECT C.regno AS regno,          SUM(F.costPerPassing) AS totalfee FROM Car C  JOIN Passing P USING(regno) JOIN TaxClass T USING(taxclas
s)
> 5000;
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
| id  | select_type | table | type   | possible_keys           | key      | key_len | ref     | rows   | Extra
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
| 1   | SIMPLE      | F    | ref    | taxclass, idx_type    | idx_type | 2       | const   | 10    | Using where; Using index; Using temporary
| 1   | SIMPLE      | T    | eq_ref | PRIMARY               | PRIMARY   | 2       | car_db.F.taxclass | 1     |
| 1   | SIMPLE      | C    | ref    | PRIMARY, taxclass, idx_regno | PRIMARY  | 21      | car_db.F.taxclass | 3593  | Using index
| 1   | SIMPLE      | S    | eq_ref | PRIMARY, idx_regno_subscription | PRIMARY  | 21      | car_db.C.regno  | 1     | Using where; Using index; Not exists
| 1   | SIMPLE      | P    | ref    | regno                 | regno    | 21      | car_db.C.regno  | 1     |
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
5 rows in set (0.000 sec)
```

d)